

**EXPERIMENT NO. 2****Aim:**

Implementation of Multiple Linear Regression, Ridge Regression, and Lasso Regression on Insurance Dataset.

**Dataset Source:**

**Dataset Name:** Insurance Premium Prediction

**Source:** Kaggle

 [Source Link:](https://www.kaggle.com/datasets/noordeen/insurance-premium-prediction/data)

<https://www.kaggle.com/datasets/noordeen/insurance-premium-prediction/data>

This dataset is publicly available and widely used for regression-based cost prediction problems.

**Dataset Description:****Overview:**

The dataset contains medical insurance information of individuals and is used to predict medical insurance charges based on demographic and lifestyle factors.

**Dataset Size:**

- Total Records: 1338
- Total Features: 7
- Problem Type: Regression
- Target Variable: Continuous (charges)

 **Feature Description**

Feature	Description
age	Age of the person
sex	Gender (male/female)
bmi	Body Mass Index
children	Number of children/dependents
smoker	Smoking status (yes/no)
region	Residential region
charges	Medical insurance cost (Target Variable)

### **Target Variable:**

charges → Medical insurance cost (continuous numeric value)

### Characteristics

- Mix of numerical and categorical data
- No missing values
- Moderate multicollinearity
- Linear relationship between BMI, smoker, and charges

## **Mathematical Formulation of the Algorithms:**

### **1. Multiple Linear Regression**

Linear Regression models the relationship between independent variables and a dependent variable.

#### **Equation:**

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n + \epsilon$$

Where:

- $y$  = Target variable (charges)
- $x_i$  = Input features
- $\beta_i$  = Coefficients
- $\epsilon$  = Error term

#### **Objective Function:**

$$MSE = \frac{1}{n} \sum (y_i - \hat{y}_i)^2$$

Minimize Mean Squared Error.

## 2. Ridge Regression (L2 Regularization)

Ridge adds penalty to reduce overfitting.

**Cost Function:**

$$J(\beta) = \frac{1}{n} \sum (y - \hat{y})^2 + \lambda \sum \beta^2$$

$\lambda$  = Regularization parameter

Penalizes large coefficients

Reduces multicollinearity impact

## 3. Lasso Regression (L1 Regularization)

Lasso performs feature selection.

**Cost Function:**

$$J(\beta) = \frac{1}{n} \sum (y - \hat{y})^2 + \lambda \sum |\beta|$$

Can shrink some coefficients to zero

Performs automatic feature selection

## Algorithm Limitations:

### 1. Linear Regression:

- Assumes linear relationship
- Sensitive to outliers
- Affected by multicollinearity
- Not suitable for non-linear datasets

## **2. Ridge Regression:**

- Does not remove irrelevant features completely
- Selection of  $\lambda$  is critical

## **3. Lasso Regression:**

- Unstable when features are highly correlated
- May eliminate useful correlated variables

## **4. Not Suitable For:**

- Highly non-linear data
- Image or text datasets without preprocessing
- Extremely small datasets

# **Methodology / Workflow:**

## **Step 1: Data Upload**

Dataset uploaded in Google Colab.

## **Step 2: Data Preprocessing**

- Automatic detection of target column
- Categorical encoding using `get_dummies()`
- Train-Test split (80-20)
- Feature scaling using `StandardScaler`

## **Step 3: Model Training**

- Linear Regression
- Ridge Regression
- Lasso Regression

#### **Step 4: Hyperparameter Tuning**

Used GridSearchCV to find optimal alpha values for Ridge and Lasso.

#### **Step 5: Model Evaluation**

Used:

- Mean Squared Error (MSE)
- R<sup>2</sup> Score

## **Performance Analysis:**

### **Evaluation Metrics Used:**

#### **1. Mean Squared Error (MSE)**

Measures prediction error.

Lower MSE → Better performance.

#### **2. R<sup>2</sup> Score:**

$$R^2 = 1 - \frac{SS_{res}}{SS_{tot}}$$

Value between 0 and 1

Closer to 1 → Better fit.

### **Observations from Output:**

- Linear Regression provides strong baseline performance.
- Ridge slightly improves generalization by shrinking coefficients.
- Lasso performs feature selection but gives similar accuracy.
- R<sup>2</sup> score around 0.75–0.80 indicates good predictive performance.

- Slight deviations in higher charge values due to outliers (e.g., smokers).

## Hyperparameter Tuning:

### 1. Ridge Regression

Parameter tuned:

$$\alpha$$

Tested values:

$$[0.01, 0.1, 1, 10, 100]$$

Used **GridSearchCV (5-fold cross-validation)**.

Best alpha selected automatically.

Impact:

- Small alpha → Less regularization
- Large alpha → Strong shrinkage
- Optimal alpha improves generalization

### 2. Lasso Regression:

Parameter tuned:

$$\lambda$$

Tested values:

$$[0.001, 0.01, 0.1, 1, 10]$$

Impact:

- Higher alpha → More coefficients shrink to zero
- Performs feature selection
- Prevents overfitting

## CODE:

```
# STEP 1: Upload File
from google.colab import files
uploaded = files.upload()

filename = list(uploaded.keys())[0]

# STEP 2: Import Libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler

# STEP 3: Load Dataset

data = pd.read_csv(filename)

print("\nColumn Names in Dataset:")
print(data.columns.tolist())

# STEP 4: Automatically Detect Target Column

# Assume last column is target
target_column = data.columns[-1]
print("\nDetected Target Column:", target_column)

y = data[target_column]
X = data.drop(target_column, axis=1)

# STEP 5: Convert Categorical to Numeric

X = pd.get_dummies(X, drop_first=True)

# STEP 6: Train-Test Split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

## # STEP 7: Feature Scaling

```
scaler = StandardScaler()  
X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)
```

## # STEP 8: Linear Regression

```
lr = LinearRegression()  
lr.fit(X_train, y_train)  
  
y_pred_lr = lr.predict(X_test)  
  
print("\n===== Linear Regression =====")  
print("MSE:", mean_squared_error(y_test, y_pred_lr))  
print("R2 Score:", r2_score(y_test, y_pred_lr))
```

## # STEP 9: Ridge Regression

```
ridge = Ridge()  
param_grid_ridge = {'alpha': [0.01, 0.1, 1, 10, 100]}  
  
grid_ridge = GridSearchCV(ridge, param_grid_ridge, cv=5)  
grid_ridge.fit(X_train, y_train)  
  
best_ridge = grid_ridge.best_estimator_  
y_pred_ridge = best_ridge.predict(X_test)  
  
print("\n===== Ridge Regression =====")  
print("Best Alpha:", grid_ridge.best_params_)  
print("MSE:", mean_squared_error(y_test, y_pred_ridge))  
print("R2 Score:", r2_score(y_test, y_pred_ridge))
```

## # STEP 10: Lasso Regression

```
lasso = Lasso(max_iter=10000)  
param_grid_lasso = {'alpha': [0.001, 0.01, 0.1, 1, 10]}  
  
grid_lasso = GridSearchCV(lasso, param_grid_lasso, cv=5)  
grid_lasso.fit(X_train, y_train)  
  
best_lasso = grid_lasso.best_estimator_  
y_pred_lasso = best_lasso.predict(X_test)
```

```
print("\n===== Lasso Regression =====")
print("Best Alpha:", grid_lasso.best_params_)
print("MSE:", mean_squared_error(y_test, y_pred_lasso))
print("R2 Score:", r2_score(y_test, y_pred_lasso))
```

## # STEP 11: Visualization

```
plt.figure(figsize=(8,6))
plt.scatter(y_test, y_pred_lr, alpha=0.5, label="Linear")
plt.scatter(y_test, y_pred_ridge, alpha=0.5, label="Ridge")
plt.scatter(y_test, y_pred_lasso, alpha=0.5, label="Lasso")

plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.title("Actual vs Predicted Comparison")
plt.legend()
plt.show()
```

## OUTPUT:

```
*** Choose Files insurance.csv
insurance.csv(text/csv) - 50264 bytes, last modified: 2/15/2026 - 100% done
Saving insurance.csv to insurance (4).csv

Column Names in Dataset:
['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'expenses']

Detected Target Column: expenses

===== Linear Regression =====
MSE: 33600065.35507784
R2 Score: 0.7835726930039905

===== Ridge Regression =====
Best Alpha: {'alpha': 10}
MSE: 33688841.98244828
R2 Score: 0.7830008582119171

===== Lasso Regression =====
Best Alpha: {'alpha': 10}
MSE: 33642353.592636935
R2 Score: 0.7833003027786798
```

Actual vs Predicted Comparison

