

EXPERIMENT NO. 3

AIM:

Decision Tree and Random Forest Classification on Insurance Dataset.

Dataset Source:

The dataset used for this experiment is the **Medical Insurance Cost Prediction Dataset**.

Source:

Kaggle – Medical Cost Personal Dataset

<https://www.kaggle.com/datasets/mirichoi0218/insurance>

(Downloaded and used locally as **insurance.csv**)

Dataset Description:

1. Dataset Overview:

- **Total Records:** 1338
- **Total Features:** 7
- **File Name:** insurance.csv
- **Type:** Tabular dataset
- **Original Problem Type:** Regression (predict insurance charges)

📌 Features Description

Feature	Description
age	Age of primary beneficiary
sex	Gender (male/female)
bmi	Body Mass Index
children	Number of children covered
smoker	Smoking status (yes/no)
region	Residential region (northeast, southeast, etc.)
charges	Medical insurance cost (continuous target)

2. Target Variable (Modified for Classification)

Since **charges** is continuous, it was converted into 3 categories using quantiles:

- 0 → Low Charges
- 1 → Medium Charges
- 2 → High Charges

This allowed us to apply **classification algorithms**.

Mathematical Formulation of the Algorithms:

A) Decision Tree:

Decision Tree splits data based on feature values using impurity measures.

♦ Gini Index Formula:

$$Gini = 1 - \sum_{i=1}^C p_i^2$$

Where:

- p_i = probability of class i
- C = number of classes

The algorithm chooses splits that minimize impurity.

Working Principle:

1. Select best feature using Gini Index.
2. Split dataset.
3. Repeat recursively.
4. Stop when:

- Maximum depth reached
- All samples belong to same class.

B) Random Forest:

Random Forest is an ensemble of multiple decision trees.

Prediction Formula:

$$\hat{y} = \text{Majority Vote of all Trees}$$

Each tree is trained on:

- Random subset of data (Bootstrap sampling)
- Random subset of features

This reduces variance and improves generalization.

Algorithm Limitations:

1. Decision Tree Limitations:

- Prone to overfitting
- Sensitive to small variations in data
- Unstable if dataset is small
- Poor performance on very complex data

2. Random Forest Limitations:

- Computationally expensive
- Less interpretable than single tree
- Requires more memory
- Slower training time

Methodology / Workflow:

Step-by-Step Workflow:

1. Load Dataset (`insurance.csv`)
2. Display basic information (shape, columns)
3. Encode categorical variables using LabelEncoder
4. Convert continuous target (`charges`) into 3 categories using `pd.qcut`
5. Split data (70% training, 30% testing)
6. Train Decision Tree
7. Train Random Forest
8. Evaluate models using:
 - Accuracy
 - Classification Report
9. Visualize:
 - Accuracy Comparison Bar Chart
 - Histogram
 - Correlation Heatmap

Performance Analysis:

a) Evaluation Metrics Used

- Accuracy
- Precision
- Recall
- F1-Score

b) Accuracy Formula:

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

Where:

- TP = True Positive
- TN = True Negative
- FP = False Positive
- FN = False Negative

c) Observations:

- Random Forest generally performs better than Decision Tree.
- Smoking status has strong impact on insurance category.
- Correlation heatmap shows:
 - High correlation between smoker and charges
 - Moderate correlation with age and BMI

Hyperparameter Tuning:

To improve performance, hyperparameter tuning was performed.

A) Decision Tree Tuning Parameters:

- `max_depth`
- `min_samples_split`
- `min_samples_leaf`

Example:

`DecisionTreeClassifier(max_depth=5, min_samples_split=4)`

Effect:

- Prevents overfitting
- Improves generalization

B) Random Forest Tuning Parameters

- `n_estimators` (number of trees)
- `max_depth`
- `min_samples_split`

Example:

`RandomForestClassifier(n_estimators=200, max_depth=10)`

Effect:

- Increasing trees improves stability
- Proper depth prevents overfitting

Impact of Tuning:

- Reduced overfitting
- Improved accuracy
- Better class balance
- More stable predictions

CODE:

```
# Decision Tree & Random Forest Classification

# Dataset: insurance.csv

import pandas as pd

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score, classification_report

from sklearn.preprocessing import LabelEncoder

#①Load Dataset

df = pd.read_csv("insurance.csv")

print("First 5 Rows:")

print(df.head())

print("\nDataset Shape:", df.shape)

print("\nColumn Names:", df.columns)

#②Convert categorical columns into numeric

for col in df.columns:

    if df[col].dtype == 'object':

        le = LabelEncoder()

        df[col] = le.fit_transform(df[col])

#③Convert LAST column into classification categories

# (because insurance dataset has continuous values)
```

```
target_column = df.columns[-1]  
df['target_category'] = pd.qcut(df[target_column],  
                                q=3,  
                                labels=[0,1,2])
```

#④ Define Features and Target

```
X = df.drop([target_column, 'target_category'], axis=1)
```

```
y = df['target_category']
```

#⑤ Train-Test Split

```
X_train, X_test, y_train, y_test = train_test_split(
```

```
    X, y, test_size=0.3, random_state=42
```

```
)
```

#⑥ Decision Tree

```
dt_model = DecisionTreeClassifier(random_state=42)
```

```
dt_model.fit(X_train, y_train)
```

```
dt_pred = dt_model.predict(X_test)
```

```
dt_accuracy = accuracy_score(y_test, dt_pred)
```

```
print("\n===== Decision Tree =====")
```

```
print("Accuracy:", dt_accuracy)
```

```
print(classification_report(y_test, dt_pred))
```

#⑦ Random Forest

```
rf_model = RandomForestClassifier(random_state=42)
```

```
rf_model.fit(X_train, y_train)
```

```
rf_pred = rf_model.predict(X_test)
```

```
rf_accuracy = accuracy_score(y_test, rf_pred)
```

```
print("\n===== Random Forest =====")  
print("Accuracy:", rf_accuracy)  
print(classification_report(y_test, rf_pred))  
# 8 Bar Chart - Accuracy Comparison  
plt.figure()  
plt.bar(["Decision Tree", "Random Forest"],  
        [dt_accuracy, rf_accuracy])  
plt.title("Model Accuracy Comparison")  
plt.xlabel("Model")  
plt.ylabel("Accuracy")  
plt.show()  
# 9 Histogram - BMI Distribution (example feature)  
plt.figure()  
plt.hist(df[df.columns[2]], bins=20)  
plt.title("Histogram of Feature")  
plt.xlabel("Value")  
plt.ylabel("Frequency")  
plt.show()  
# 10 Correlation Heatmap  
plt.figure()  
correlation_matrix = df.drop('target_category', axis=1).corr()  
plt.imshow(correlation_matrix)  
plt.title("Correlation Heatmap")  
plt.xticks(range(len(correlation_matrix.columns)),  
           correlation_matrix.columns,
```

```

rotation=90)

plt.yticks(range(len(correlation_matrix.columns)),
           correlation_matrix.columns)

plt.colorbar()

plt.show()

print("\nExperiment Completed Successfully!")

```

OUTPUT:

1.

```

First 5 Rows:
   age   sex   bmi  children smoker      region  expenses
0   19 female  27.9        0    yes southwest  16884.92
1   18 male   33.8        1    no southeast  1725.55
2   28 male   33.0        3    no southeast  4449.46
3   33 male   22.7        0    no northwest 21984.47
4   32 male   28.9        0    no northwest  3866.86

Dataset Shape: (1338, 7)

Column Names: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'expenses'], dtype='object')

===== Decision Tree =====
Accuracy: 0.8532338308457711
      precision    recall  f1-score   support
          0         0.90     0.93     0.91      143
          1         0.84     0.80     0.82      132
          2         0.81     0.82     0.82      127

      accuracy                           0.85      402
      macro avg       0.85     0.85     0.85      402
  weighted avg       0.85     0.85     0.85      402

```

2.


```

===== Random Forest =====
Accuracy: 0.8930348258706468
      precision    recall  f1-score   support
          0         0.88     0.94     0.91      143
          1         0.87     0.92     0.89      132
          2         0.94     0.81     0.87      127


      accuracy                           0.89      402
      macro avg       0.90     0.89     0.89      402
  weighted avg       0.90     0.89     0.89      402

```


3.



4.



5.



Experiment Completed Successfully!